

# ridgeRegressionAndLASSO

October 28, 2019

## 1 Ridge Regression

Our goal is to learn about Ridge Regression. This technique is closely related to standard linear regression, but adds an extra twist. The twist involves including an additional term in our minimization of sums of squares that encourages coefficients to shrink.

The idea behind ridge regression is that influencing coefficients to shrink towards zero will also discourage overfitting to our training data. If we can prevent overfitting then we can better generalize our model to yet uncollected data.

### 1.1 Mechanics

We saw in previous lectures that finding optimal coefficients in a linear regression is the same as minimizing the sum squares error. More concretely, given a set of observations  $(x, y)_1, (x, y)_2, \dots, (x, y)_N$  we can write our probabilistic model as

$$p(y|x) \sim N(\beta_0 + \beta'x, \sigma^2)$$

where  $\beta_0$  is an intercept.

The optimal  $\beta$  parameters are found by minimizing the following function of  $\beta$

$$\min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \beta'x_i)^2 \right\}$$

Ridge regression adds an additional term to the above optimization problem

$$\min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \beta'x_i)^2 \right\} + \lambda \sum_{m=1}^M \beta_m^2$$

The  $\lambda$  parameter is a choice on part of the investigator. Typically,  $\lambda$  is chosen by cross-validation.

This additional term penalizes coefficients related to covariates. Ridge regression does **not** penalize the size of the intercept.

## 1.2 Optimal $\beta$

We can solve our optimization problem by taking partial derivatives and finding the  $\beta$  that zeros out these derivatives.

$$f(\beta) = \sum_{i=1}^N (y_i - \beta_0 - \beta' x_i)^2 + \lambda \sum_{m=1}^M \beta_m^2$$

The partial derivative with respect to  $\beta_m$  is

$$\frac{df}{d\beta_m} = \sum_{i=1}^N -2x_{im} (y_i - \beta_0 - \beta' x_i) + 2\lambda\beta_m \quad (1)$$

We set the derivative equal to zero and solve for  $\beta_m$ .

$$\sum_{i=1}^N x_{im} (y_i - \beta_0 - \beta' x_i) - \lambda\beta_m = 0 \quad (2)$$

$$\sum_{i=1}^N x_{im} y_i - \sum_{i=1}^N x_{im} (\beta_0 + \beta' x_i) - \lambda\beta_m = 0 \quad (3)$$

$$\sum_{i=1}^N x_{im} y_i = \sum_{i=1}^N x_{im} (\beta_0 + \beta' x_i) + \lambda\beta_m \quad (4)$$

$$(5)$$

$$x'_m y = x'_m (X + \lambda 1_m) \beta \quad (6)$$

where  $1_m$  is a vector of zeros everywhere and a 1 in the  $m^{th}$  position. The set of all optimal betas is then

$$X' y = (X' X + \lambda I) \beta$$

and the optimal  $\beta$  for ridge regression is

$$\beta_{RR} = (X' X + \lambda I)^{-1} X' y$$

This vector of  $\beta$ s does not include the intercept.

## 2 LASSO

LASSO works similarly to ridge regression. The goal again is to optimize the typical SSE with an additional term

$$\min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \beta' x_i)^2 \right\} + \lambda \sum_{m=1}^M |\beta_m|$$

The additional term is the sum of the absolute value of  $\beta$ s, not including the intercept.

Though a subtle difference, we can not solve for the optimal  $\beta$ s exactly. A numerical solver capable of convex optimization can provide an estimate of the optimal  $\beta$ s for lasso.

The absolute value constraint put on the sum of coefficients is stronger than the squared coefficients in ridge regression. This additional pressure to shrink estimates serves as a natural way to, among many candidate coefficients, find a subset that best predict the target variable  $y$ .

### 2.1 Example data

The example data set is a collection of information about baseball players. The data contains, for each baseball player, information about their performance and how much money they make.

Our goal will be to understand the relationship between a baseball player's performance and their salary.

To generalize from this subset to other baseball players we did not collect data on, ridge regression will be used. Ridge regression will penalize over emphasizing specific covariates that could cause us to overfit to our training data. We can also fit a LASSO model to our training data. LASSO will shrink covariate estimates to zero if they do not substantially contribute to reducing the SSE. LASSO aims to create a more parsimonious model, and by doing so, create a model that can better generalize from training to testing data.

```
[56]: library(ISLR)
library(glmnet)
library(dplyr)
library(tidyr)

# Hitters dataset
Hitters = na.omit(Hitters)
print(head(Hitters))

# Take covariates except Salary
x = model.matrix(Salary~., Hitters)[,-1]

# Take salary variable
y = Hitters$Salary

# Ridge regression can be fit with the glmnet function.
```

```

# This function takes a matrix of covariates, the target variable.
# Setting alpha=0 specifies Ridge regression and Lambda corresponds the the
  ↳ amount of covariate "shrinkage"
grid = 10^seq(10, -2, length = 100)
ridge_mod = glmnet(x, y, alpha = 0, lambda = grid)

C = coef(ridge_mod)
for (k in 1:nrow(C)){
  plot( log(grid), C[k,], type='l', ylim=c(-100,500))
  par(new=TRUE)
}

```


	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun
-Alan Ashby	315	81	7	24	38	39	14	3449	835	69
-Alvin Davis	479	130	18	66	72	76	3	1624	457	63
-Andre Dawson	496	141	20	65	78	37	11	5628	1575	225
-Andres Galarraga	321	87	10	39	42	30	2	396	101	12
-Alfredo Griffin	594	169	4	74	51	35	11	4408	1133	19
-Al Newman	185	37	1	23	8	21	2	214	42	1

	CRuns	CRBI	CWalks	League	Division	PutOuts	Assists	Errors
-Alan Ashby	321	414	375	N	W	632	43	10
-Alvin Davis	224	266	263	A	W	880	82	14
-Andre Dawson	828	838	354	N	E	200	11	3
-Andres Galarraga	48	46	33	N	E	805	40	4
-Alfredo Griffin	501	336	194	A	W	282	421	25
-Al Newman	30	9	24	N	E	76	127	7

	Salary	NewLeague
-Alan Ashby	475.0	N
-Alvin Davis	480.0	A
-Andre Dawson	500.0	N
-Andres Galarraga	91.5	N
-Alfredo Griffin	750.0	A
-Al Newman	70.0	A



```
[63]: cat("Lambda")
ridge_mod$lambda[60]
coef(ridge_mod)[,60]

cat("SSE")
sqrt(sum(coef(ridge_mod)[-1,60]^2))

cat("Lambda")
ridge_mod$lambda[50]
coef(ridge_mod)[,50]

cat("SSE")
sqrt(sum(coef(ridge_mod)[-1,50]^2))
```

Lambda

705.480231071865

**(Intercept)** 54.3251995018372 **AtBat** 0.112111145878249 **Hits** 0.656224085323628 **HmRun** 1.17980909638777 **Runs** 0.937697128927054 **RBI** 0.847185458771521 **Walks** 1.31987948048781 **Years** 2.59640424574253 **CAtBat** 0.0108341254432856 **CHits** 0.0467455700054452 **CHmRun** 0.337773183143353 **CRuns** 0.0935552830000676 **CRBI** 0.0978040232271687 **CWalks** 0.0718961166304866 **LeagueN** 13.6837019095343 **DivisionW** -54.658777504592 **PutOuts** 0.118522894134745 **Assists** 0.01606037317599 **Errors** -0.703586547290985 **NewLeagueN** 8.61181213448926

SSE

57.110014262533

Lambda

11497.5699539774

**(Intercept)** 407.356050200416 **AtBat** 0.0369571817501359 **Hits** 0.138180343807892 **HmRun** 0.524629975886911 **Runs** 0.230701522621179 **RBI** 0.239841458504058 **Walks** 0.289618741049884 **Years** 1.10770292908555 **CAtBat** 0.00313181522151328 **CHits** 0.0116536373557531 **CHmRun** 0.0875456697555949 **CRuns** 0.0233798823693758 **CRBI** 0.0241383203685686 **CWalks** 0.0250154205993732 **LeagueN** 0.0850281135625444 **DivisionW** -6.21544097273146 **PutOuts** 0.0164825767604547 **Assists** 0.00261298804528183 **Errors** -0.0205026903654579 **NewLeagueN** 0.301433531372699

SSE

6.36061242142791

```
[67]: # The same process used for Ridge Regression can be used to fit LASSO.
# The only change is specifying alpha=1.
grid = 10^seq(10, -2, length = 100)
ridge_mod = glmnet(x, y, alpha = 1, lambda = grid)

cat("Lambda")
ridge_mod$lambda[90]
coef(ridge_mod)[,90]
sqrt(sum(coef(ridge_mod)[-1,90]^2))

cat("Lambda")
ridge_mod$lambda[80]
coef(ridge_mod)[,80]
sqrt(sum(coef(ridge_mod)[-1,80]^2))

cat("Lambda")
ridge_mod$lambda[70]
coef(ridge_mod)[,70]
sqrt(sum(coef(ridge_mod)[-1,70]^2))
```

Lambda

0.162975083462064

**(Intercept)** 162.471571163125 **AtBat** -2.00725495341728 **Hits** 7.46834345075141 **HmRun**  
3.50180353190197 **Runs** -2.15880439279739 **RBI** -0.741855229071078 **Walks** 6.15950199852769  
**Years** -4.69367835398105 **CAtBat** -0.140363573272626 **CHits** 0.0843854865595957 **CHmRun**  
-0.00859358349740949 **CRuns** 1.39671552949042 **CRBI** 0.724102431969445 **CWalks**  
-0.810053300830832 **LeagueN** 59.7052960522582 **DivisionW** -116.71196728165 **PutOuts**  
0.283100309617026 **Assists** 0.354242405196464 **Errors** -3.25345107822638 **NewLeagueN**  
-22.0755201272487

133.510043789335

Lambda

2.65608778294668

**(Intercept)** 124.089487297287 **AtBat** -1.56009839061161 **Hits** 5.69316850453147 **HmRun** 0 **Runs** 0  
**RBI** 0 **Walks** 4.7505395162242 **Years** -9.51802414366672 **CAtBat** 0 **CHits** 0 **CHmRun**  
0.519161053589954 **CRuns** 0.660407436641157 **CRBI** 0.391541539043815 **CWalks**  
-0.53268682707415 **LeagueN** 32.112549338136 **DivisionW** -119.25835400136 **PutOuts**  
0.272620661250992 **Assists** 0.174816439652482 **Errors** -2.05674316297772 **NewLeagueN** 0

124.125964686996

Lambda

43.2876128108306

**(Intercept)** 74.7044627184386 **AtBat** 0 **Hits** 1.64328660554985 **HmRun** 0 **Runs** 0 **RBI** 0 **Walks**  
1.89863540344793 **Years** 0 **CAtBat** 0 **CHits** 0 **CHmRun** 0 **CRuns** 0.183705603033279 **CRBI**  
0.374319642602473 **CWalks** 0 **LeagueN** 0 **DivisionW** -55.4369608623154 **PutOuts**  
0.151945324305305 **Assists** 0 **Errors** 0 **NewLeagueN** 0

55.4955744833595